

## ✓ Deep Learning - Training a Simple Convolution Neural Network Model

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import CategoricalCrossentropy
import matplotlib.pyplot as plt

# Step 1: Define image dimensions, batch size, and directory paths
image_height = 100
image_width = 100
batch_size = 32

train_dir = '/content/drive/MyDrive/TrainingImages/'
test_dir = '/content/drive/MyDrive/TestingImages/'

# Step 2: Data generators for training and testing
train_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)
test_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)

train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(image_height, image_width),
    batch_size=batch_size,
    class_mode='categorical'
)

test_generator = test_datagen.flow_from_directory(
    test_dir,
    target_size=(image_height, image_width),
    batch_size=batch_size,
    class_mode='categorical'
)

# Step 3: Define the model architecture
model = Sequential([
    Conv2D(8, (3, 3), activation='relu', input_shape=(image_height, image_width, 3)),
    Conv2D(8, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(16, activation='relu'),
    Dense(16, activation='relu'),
    Dense(10, activation='softmax')
])

# Step 4: Compile the model
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Step 5: Train the model
```

```
history = model.fit(train_generator,
                    epochs=20,
                    validation_data=test_generator)

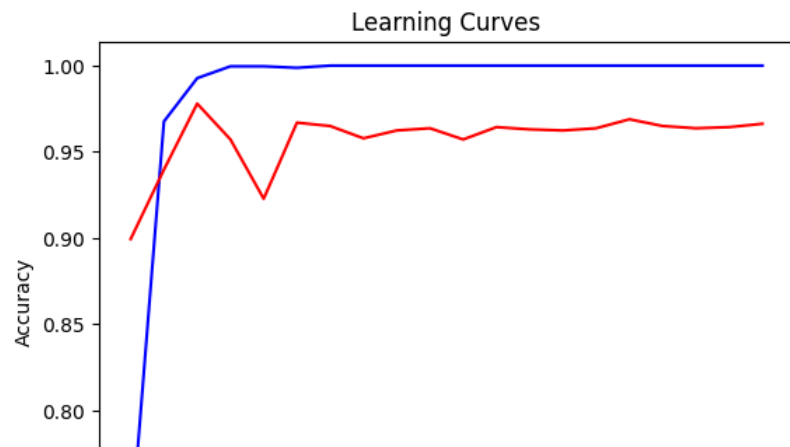
# Step 6(a): Plot learning curves
plt.plot(history.history['accuracy'], label='Training Accuracy', color='blue')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy', color='red')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Learning Curves')
plt.legend()
plt.show()
```


Found 4665 images belonging to 10 classes.  
 Found 1542 images belonging to 10 classes.

```

Epoch 1/20
146/146 [=====] - 2997s 21s/step - loss: 0.8114 - accuracy: 0.7269 - val_loss: 0.2829 - val_accuracy: 0.8995
Epoch 2/20
146/146 [=====] - 54s 368ms/step - loss: 0.0866 - accuracy: 0.9676 - val_loss: 0.1882 - val_accuracy: 0.9397
Epoch 3/20
146/146 [=====] - 54s 366ms/step - loss: 0.0313 - accuracy: 0.9927 - val_loss: 0.0659 - val_accuracy: 0.9780
Epoch 4/20
146/146 [=====] - 54s 368ms/step - loss: 0.0067 - accuracy: 0.9996 - val_loss: 0.0975 - val_accuracy: 0.9572
Epoch 5/20
146/146 [=====] - 51s 347ms/step - loss: 0.0035 - accuracy: 0.9996 - val_loss: 0.2096 - val_accuracy: 0.9228
Epoch 6/20
146/146 [=====] - 52s 354ms/step - loss: 0.0045 - accuracy: 0.9987 - val_loss: 0.0972 - val_accuracy: 0.9669
Epoch 7/20
146/146 [=====] - 53s 366ms/step - loss: 0.0011 - accuracy: 1.0000 - val_loss: 0.0683 - val_accuracy: 0.9650
Epoch 8/20
146/146 [=====] - 54s 366ms/step - loss: 6.2037e-04 - accuracy: 1.0000 - val_loss: 0.0864 - val_accuracy: 0.9578
Epoch 9/20
146/146 [=====] - 52s 352ms/step - loss: 4.3142e-04 - accuracy: 1.0000 - val_loss: 0.0817 - val_accuracy: 0.9624
Epoch 10/20
146/146 [=====] - 52s 358ms/step - loss: 2.9534e-04 - accuracy: 1.0000 - val_loss: 0.0742 - val_accuracy: 0.9637
Epoch 11/20
146/146 [=====] - 50s 337ms/step - loss: 2.7244e-04 - accuracy: 1.0000 - val_loss: 0.0932 - val_accuracy: 0.9572
Epoch 12/20
146/146 [=====] - 49s 333ms/step - loss: 2.3860e-04 - accuracy: 1.0000 - val_loss: 0.0728 - val_accuracy: 0.9643
Epoch 13/20
146/146 [=====] - 52s 357ms/step - loss: 1.6473e-04 - accuracy: 1.0000 - val_loss: 0.0791 - val_accuracy: 0.9630
Epoch 14/20
146/146 [=====] - 49s 334ms/step - loss: 1.4142e-04 - accuracy: 1.0000 - val_loss: 0.0815 - val_accuracy: 0.9624
Epoch 15/20
146/146 [=====] - 54s 372ms/step - loss: 1.1591e-04 - accuracy: 1.0000 - val_loss: 0.0852 - val_accuracy: 0.9637
Epoch 16/20
146/146 [=====] - 52s 359ms/step - loss: 1.0319e-04 - accuracy: 1.0000 - val_loss: 0.0725 - val_accuracy: 0.9689
Epoch 17/20
146/146 [=====] - 52s 357ms/step - loss: 9.3892e-05 - accuracy: 1.0000 - val_loss: 0.0751 - val_accuracy: 0.9650
Epoch 18/20
146/146 [=====] - 49s 338ms/step - loss: 7.7498e-05 - accuracy: 1.0000 - val_loss: 0.0832 - val_accuracy: 0.9637
Epoch 19/20
146/146 [=====] - 52s 354ms/step - loss: 7.0125e-05 - accuracy: 1.0000 - val_loss: 0.0868 - val_accuracy: 0.9643
Epoch 20/20
146/146 [=====] - 49s 331ms/step - loss: 6.0045e-05 - accuracy: 1.0000 - val_loss: 0.0780 - val_accuracy: 0.9663

```





## ✓ Transfer Learning via Feature Extraction

```
import torch
import torchvision
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
import numpy as np
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
from sklearn.decomposition import PCA
from sklearn.model_selection import train_test_split

# Step 1: Define configuration variables
image_height = 224 # Modified image height
image_width = 224 # Modified image width
batch_size = 32
num_classes = 10
num_epochs = 10
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Step 2: Load the dataset using ImageDataLoader and organize it on disk
train_dir = "/content/drive/MyDrive/TrainingImages/"
test_dir = "/content/drive/MyDrive/TestingImages/"

train_dataset = torchvision.datasets.ImageFolder(
    root=train_dir,
    transform=torchvision.transforms.Compose([
        torchvision.transforms.Resize((image_height, image_width)),
        torchvision.transforms.ToTensor(),
    ])
)

test_dataset = torchvision.datasets.ImageFolder(
    root=test_dir,
    transform=torchvision.transforms.Compose([
        torchvision.transforms.Resize((image_height, image_width)),
        torchvision.transforms.ToTensor(),
    ])
)

# Split the training dataset into training and validation sets
train_data, val_data = train_test_split(train_dataset, test_size=0.2, random_state=42)
```

```
train_data, val_data = train_test_split(train_dataset, test_size=0.2, random_state=12,
```

```
# Define data loaders for training, validation, and testing
```

```
train_loader = torch.utils.data.DataLoader(
    train_data,
    batch_size=batch_size,
    shuffle=True
)
```

```
val_loader = torch.utils.data.DataLoader(
    val_data,
    batch_size=batch_size,
    shuffle=False
)
```

```
test_loader = torch.utils.data.DataLoader(
    test_dataset,
    batch_size=batch_size,
    shuffle=False
)
```

```
# Define a function to calculate accuracy
```

```
def calculate_accuracy(loader, model):
    model.eval()
    correct = 0
    total = 0
    with torch.no_grad():
        for images, labels in loader:
            images, labels = images.to(device), labels.to(device)
            outputs = model(images)
            _, predicted = torch.max(outputs.data, 1)
            total += labels.size(0)
            correct += (predicted == labels).sum().item()
    return correct / total
```

```
# Define the model, criterion, and optimizer
```

```
resnet18 = torchvision.models.resnet18(pretrained=True)
for param in resnet18.parameters():
    param.requires_grad = False
```

```
num_fts = resnet18.fc.in_features
resnet18.fc = nn.Linear(num_fts, num_classes)
resnet18 = resnet18.to(device)
```

```
criterion = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(resnet18.parameters(), lr=0.001)
```

```
# Lists to store training and validation accuracies
```

```
train_accuracies, val_accuracies = [], []
```

```
# Train the model
```

```
for epoch in range(num_epochs):
    running_loss = 0.0
    correct_train, total_train = 0, 0
    resnet18.train()
    for images, labels in train_loader:
        images, labels = images.to(device), labels.to(device)
```

```

images, labels = images.to(device), labels.to(device)
optimizer.zero_grad()
outputs = resnet18(images)
loss = criterion(outputs, labels)
loss.backward()
optimizer.step()
running_loss += loss.item()

# Calculate training accuracy
_, predicted = torch.max(outputs.data, 1)
total_train += labels.size(0)
correct_train += (predicted == labels).sum().item()

# Calculate training accuracy
train_accuracy = correct_train / total_train
train_accuracies.append(train_accuracy)

# Calculate validation accuracy
val_accuracy = calculate_accuracy(val_loader, resnet18)
val_accuracies.append(val_accuracy)

print("Epoch: {}/{}.. ".format(epoch+1, num_epochs),
      "Training Loss: {:.3f}.. ".format(running_loss/len(train_loader)),
      "Training Accuracy: {:.3f}.. ".format(train_accuracy),
      "Validation Accuracy: {:.3f}".format(val_accuracy))

# Plot learning curves
epochs = range(1, num_epochs + 1)
plt.plot(epochs, train_accuracies, label='Training Accuracy', color='blue')
plt.plot(epochs, val_accuracies, label='Validation Accuracy', color='red')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.title('Learning Curves')
plt.legend()
plt.show()

# Step 6: Extract features from ResNet18 for training and test datasets
class FeatureExtractor:
    def __init__(self, model):
        self.model = model
        self.features = None
        self.hook = self.model.layer4.register_forward_hook(self.hook_fn)

    def hook_fn(self, module, input, output):
        self.features = output

    def extract(self, x):
        self.model(x)
        return self.features

train_feature_extractor = FeatureExtractor(resnet18)
test_feature_extractor = FeatureExtractor(resnet18)

train_features = []
train_labels = []

```

```
test_features = []
test_labels = []

# Extract features and labels from the training dataset
for images, labels in train_loader:
    features_batch = train_feature_extractor.extract(images.to(device))
    train_features.append(features_batch.detach().cpu().numpy())
    train_labels.append(labels.numpy())

# Extract features and labels from the test dataset
for images, labels in test_loader:
    features_batch = test_feature_extractor.extract(images.to(device))
    test_features.append(features_batch.detach().cpu().numpy())
    test_labels.append(labels.numpy())

train_features = np.concatenate(train_features)
train_labels = np.concatenate(train_labels)
test_features = np.concatenate(test_features)
test_labels = np.concatenate(test_labels)

# Step 7: Perform PCA dimensionality reduction
pca = PCA(n_components=2)

# Flatten the features
train_features_flattened = train_features.reshape(train_features.shape[0], -1)
test_features_flattened = test_features.reshape(test_features.shape[0], -1)

# Apply PCA transformation
train_features_pca = pca.fit_transform(train_features_flattened)
test_features_pca = pca.transform(test_features_flattened)

# Step 8: Train SVM prediction model using extracted features
svm_model = SVC(kernel='rbf', C=10)
svm_model.fit(train_features_flattened, train_labels)

# Step 9: Evaluate SVM model
train_predictions = svm_model.predict(train_features_flattened)
test_predictions = svm_model.predict(test_features_flattened)

train_accuracy = accuracy_score(train_labels, train_predictions)
test_accuracy = accuracy_score(test_labels, test_predictions)

print("SVM Training Accuracy", train_accuracy)
print("SVM Test Accuracy", test_accuracy)
```

```

/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed :
warnings.warn(
/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated
warnings.warn(msg)

```


Downloading: "<https://download.pytorch.org/models/resnet18-f37072fd.pth>" to /root/.cache/torch/hub/checkpoints/resnet18-f37072fd.pth

100%|██████████| 44.7M/44.7M [00:00<00:00, 116MB/s]

```

Epoch: 1/10.. Training Loss: 0.493.. Training Accuracy: 0.919.. Validation Accuracy: 1.000
Epoch: 2/10.. Training Loss: 0.058.. Training Accuracy: 0.999.. Validation Accuracy: 1.000
Epoch: 3/10.. Training Loss: 0.028.. Training Accuracy: 1.000.. Validation Accuracy: 1.000
Epoch: 4/10.. Training Loss: 0.018.. Training Accuracy: 1.000.. Validation Accuracy: 1.000
Epoch: 5/10.. Training Loss: 0.013.. Training Accuracy: 1.000.. Validation Accuracy: 1.000
Epoch: 6/10.. Training Loss: 0.010.. Training Accuracy: 1.000.. Validation Accuracy: 1.000
Epoch: 7/10.. Training Loss: 0.007.. Training Accuracy: 1.000.. Validation Accuracy: 1.000
Epoch: 8/10.. Training Loss: 0.007.. Training Accuracy: 1.000.. Validation Accuracy: 1.000
Epoch: 9/10.. Training Loss: 0.005.. Training Accuracy: 1.000.. Validation Accuracy: 1.000
Epoch: 10/10.. Training Loss: 0.005.. Training Accuracy: 1.000.. Validation Accuracy: 1.000

```



SVM Training Accuracy 1.0  
SVM Test Accuracy 1.0

## ✓ Transfer Learning via Fine-Tuning



```
import torch
import torchvision
import torch.nn as nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
import numpy as np
from sklearn.model_selection import train_test_split

# Step 1: Define configuration variables
image_height = 224 # Modified image height
image_width = 224 # Modified image width
batch_size = 32
num_classes = 10
num_epochs = 10
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")

# Step 2: Load the dataset using ImageDataLoader and organize it on disk
train_dir = "/content/drive/MyDrive/TrainingImages/"
test_dir = "/content/drive/MyDrive/TestingImages/"

train_dataset = torchvision.datasets.ImageFolder(
    root=train_dir,
    transform=torchvision.transforms.Compose([
        torchvision.transforms.Resize((image_height, image_width)),
        torchvision.transforms.ToTensor(),
    ])
)

test_dataset = torchvision.datasets.ImageFolder(
    root=test_dir,
    transform=torchvision.transforms.Compose([
        torchvision.transforms.Resize((image_height, image_width)),
        torchvision.transforms.ToTensor(),
    ])
)

# Split the training dataset into training and validation sets
train_data, val_data = train_test_split(train_dataset, test_size=0.2, random_state=42)

train_loader = torch.utils.data.DataLoader(
    train_data,
    batch_size=batch_size,
    shuffle=True
)

test_loader = torch.utils.data.DataLoader(
    test_dataset,
    batch_size=batch_size,
    shuffle=False
)

# Step 3: Build the model (fine-tuning)
resnet18 = torchvision.models.resnet18(weights=None)
num_fts = resnet18.fc.in_features
resnet18.fc = nn.Linear(num_fts, num_classes).to(device)
```



```
resnet18 = resnet18.to(device)
```